

# Neural Networks (2014)

## Practical Assignment III: Learning by gradient descent

### Stochastic gradient descent

The aim of this problem is to get acquainted with gradient descent based training in practice and do some *hands on* experiments. Take the actual assignment as a starting point for further exploration and self-study.

We consider a simple feedforward neural network with real-valued output

$$\sigma(\xi) = (\tanh[\mathbf{w}_1 \cdot \xi] + \tanh[\mathbf{w}_2 \cdot \xi])$$

where  $\xi \in \mathbb{R}^N$  represents an input vector and  $\mathbf{w}_1$  and  $\mathbf{w}_2$  are the  $N$ -dim. vectors of adaptive input-to-hidden weights. The fixed hidden-to-output relation is given as the sum of the hidden states (*soft committee machine*).

#### a) Stochastic gradient descent

Formulate and implement a stochastic gradient descent procedure w.r.t. the weight vectors  $\mathbf{w}_1$  and  $\mathbf{w}_2$ , which aims at the minimization of the cost function

$$E = \frac{1}{P} \frac{1}{2} \sum_{\mu=1}^P (\sigma(\xi^\mu) - \tau(\xi^\mu))^2$$

for a given data set  $\mathcal{D} = \{\xi^\mu, \tau(\xi^\mu)\}_{\mu=1}^P$  with continuous training labels  $\tau(\xi^\mu) \in \mathbb{R}$ .

In each learning step, select one of the  $P$  examples, say  $\xi^\nu$ , randomly with equal probability and use the gradient with respect to its contribution  $e^\nu = (\sigma(\xi^\nu) - \tau(\xi^\nu))^2/2$ , only (stochastic gradient descent):

$$\begin{aligned}\mathbf{w}_1 &\leftarrow \mathbf{w}_1 - \eta \nabla_{\mathbf{w}_1} e^\nu \\ \mathbf{w}_2 &\leftarrow \mathbf{w}_2 - \eta \nabla_{\mathbf{w}_2} e^\nu\end{aligned}$$

where  $\nabla_{\mathbf{w}_j}$  denotes the gradient with respect to  $\mathbf{w}_j$ .

Perform  $t_{max} \cdot P$  single training steps, where  $t \leq t_{max}$  measures the *training time* in units of  $P$  single updates.

Initialize the weights as independent random vectors with  $|\mathbf{w}_1|^2 = 1$  and  $|\mathbf{w}_2|^2 = 1$ . You can work with a constant learning rate. A reasonable choice of the learning rate should be  $\eta = 0.05$ , but you may want to use different values, depending on the observed performance.

#### b) A regression problem

In Nestor you will find the file `data3.mat` which you should import into Matlab (`load data3.mat`). It provides a  $50 \times 5000$ -dim. array `xi` corresponding to 5000 input vectors (dimension  $N = 50$ ) and a 5000-dim. vector `tau` corresponding to the target values.

Consider (at least) the first  $P = 100$  examples as the training set.

In the course of stochastic gradient descent training, measure the cost function  $E$  and plot it vs. the time  $t$  as defined above.

In addition, evaluate the quantity

$$E_{test} = \frac{1}{Q} \frac{1}{2} \sum_{\rho=P+1}^{P+Q} (\sigma(\xi^\rho) - \tau(\xi^\rho))^2$$

which corresponds to the *test* or *generalization error* in terms of quadratic deviation from the target function for  $Q$  test examples, set  $Q = 100$  or larger.

Plot and compare the evolution of  $E$  and  $E_{test}$  with the training time  $t$ . You should consider training times  $t_{max}$  after which the errors seem to have become roughly constant (apart from

fluctuations). Display the obtained, final weight vectors as bar graphs.

Hand in at least the following:

- A brief description of the problem, including the update equations according to stochastic gradient descent
- The curves  $E(t)$  and  $E_{test}(t)$  corresponding to the above specified (or similar) parameters
- Bar graphs displaying the two weight vectors after  $t_{max}$

#### Remarks:

- You can follow the supplementary material (grad-example.pdf) in Nestor for the calculation of derivatives. Note, however, that here hidden-to-output weights and the gain parameters are fixed.
- If you want to solve the problem using some other programming language you can export the arrays in text or csv format from within matlab using the commands `save` or `csvwrite`, for instance. See the matlab documentation for syntax and options.
- You could also consider the concatenated vector  $\underline{W} = [\mathbf{w}_1, \mathbf{w}_2]$  and use the gradient with respect to  $\underline{W}$ , but since  $\nabla_{\underline{W}} = [\nabla_{\mathbf{w}_1}, \nabla_{\mathbf{w}_2}]$  this is completely equivalent to the above.
- Compute  $E$  and  $E_{test}$  after  $P$  single randomized steps, not after each individual update.
- Of course, your results will be more reliable if you repeat the training process over several runs from random initializations and take an average of  $E(t)$  and  $E_{test}(t)$  over these runs. However, this is not obligatory and depends on your patience and available CPU time.

---

#### Possible bonus problems

- consider smaller and larger values of  $P$ , e.g. a selection from  $P = 20, 50, 200, 500, 1000, 2000$  for the training process. How do the final training and test errors depend on  $P$ ? Make sure that the  $Q$  test examples are not used in the training, of course.
- study systematically the influence of the learning rate  $\eta$ . Potentially consider a time dependent rate as discussed in class.
- Can you observe plateau states? If so, display the corresponding weight vectors and compare with the final ones after leaving the plateau.
- consider a student network with  $K > 2$  for the same data set. Can you observe over-fitting?
- ...